

MIT OpenCourseWare
<http://ocw.mit.edu>

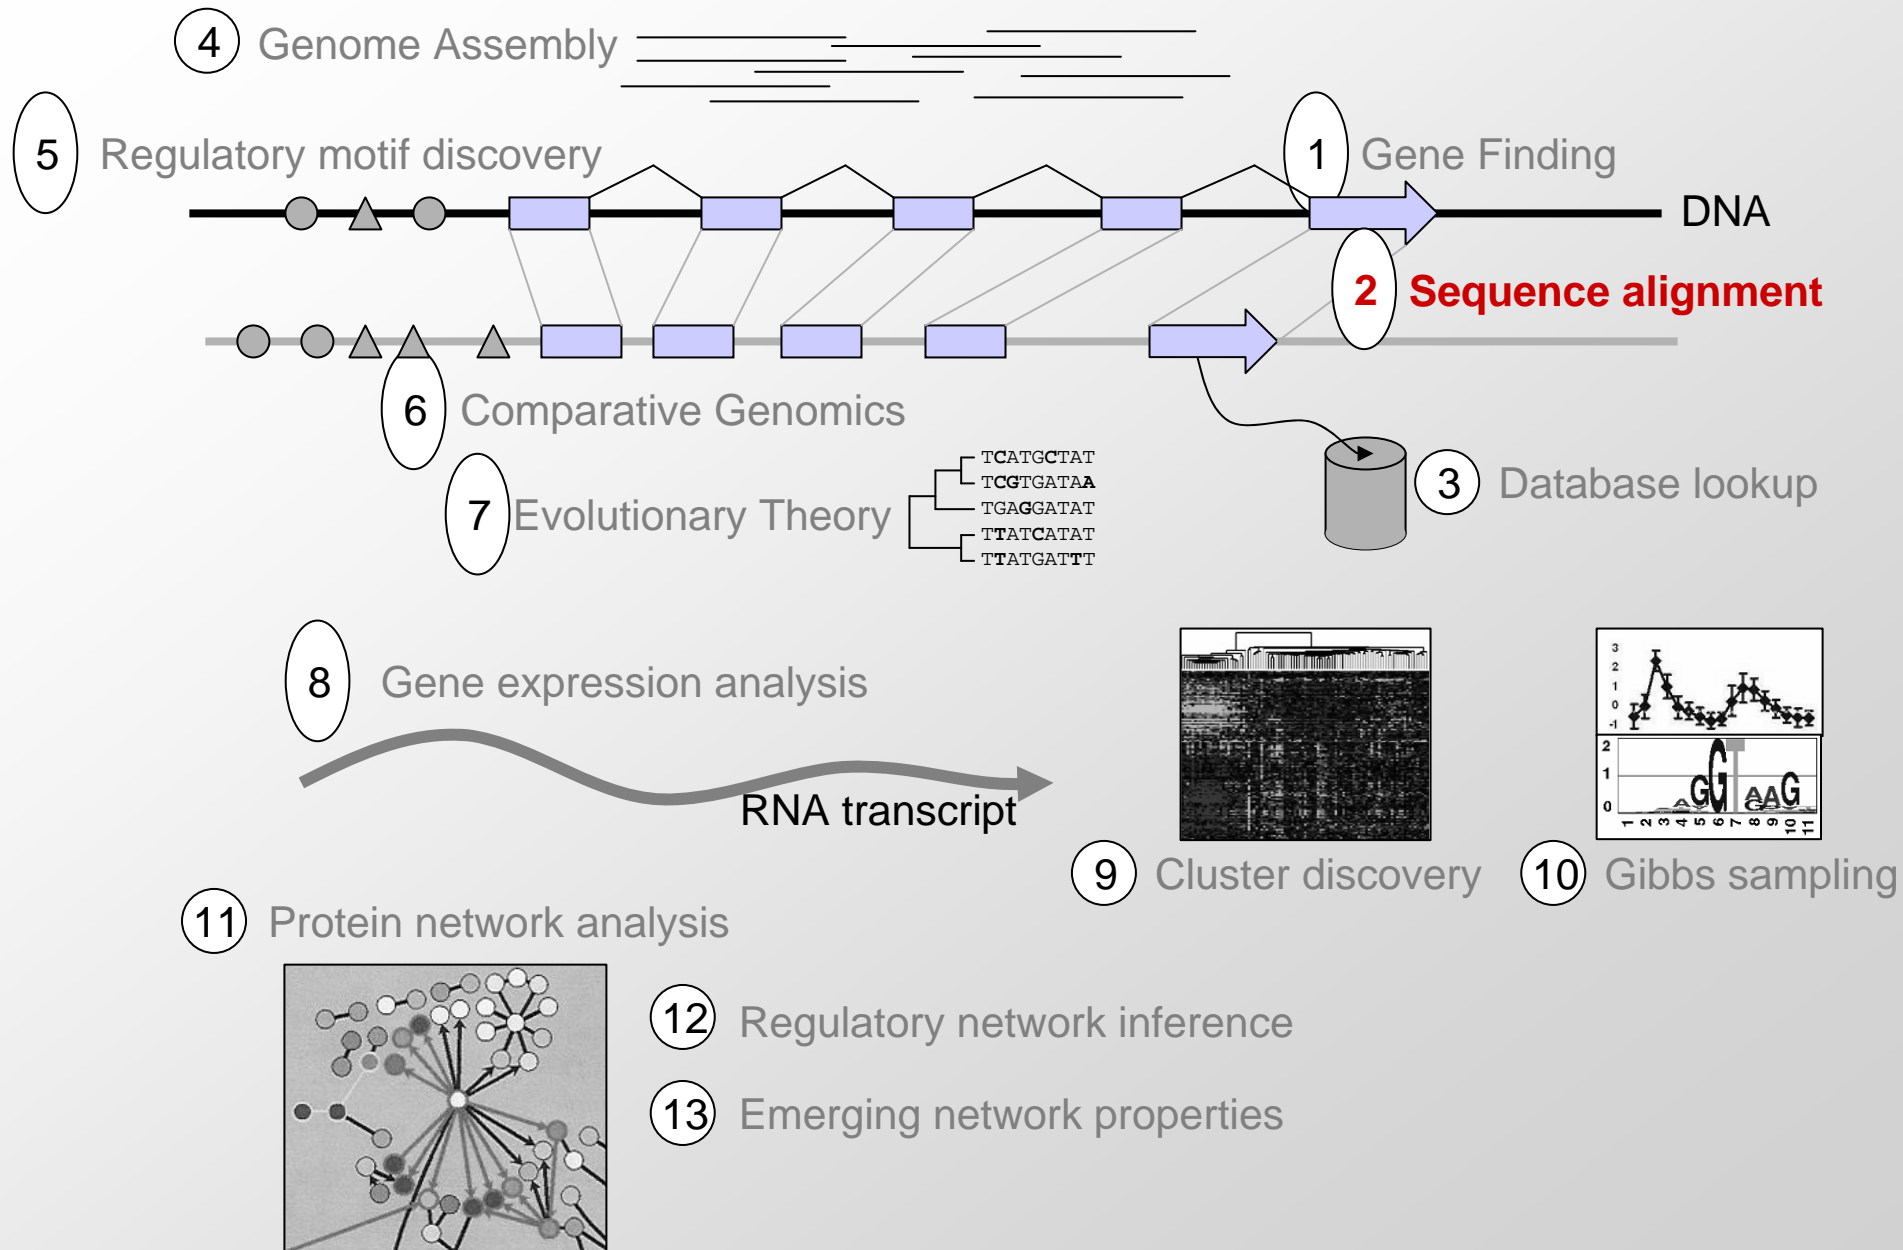
6.047 / 6.878 Computational Biology: Genomes, Networks, Evolution
Fall 2008

For information about citing these materials or our Terms of Use, visit: <http://ocw.mit.edu/terms>.

Sequence Alignment and Dynamic Programming

Tue Sept 9, 2008


Challenges in Computational Biology



Reminder: Last lecture / recitation

- Schedule for the term
 - ‘Foundations’ till midterm
 - ‘Frontiers’ lead to final project
 - Duality: basic problems / fundamental techniques
- Biology introduction
 - DNA, RNA, protein, transcription, translation
 - Why computational biology
- Today: Comparative genomics is everywhere!
 - Problem set 1: dating vertebrate whole-genome duplication
 - Problem set 2: discover genes using their conservation properties
 - Problem set 3: discover all motifs across entire yeast genome
 - Problem set 4: reversing human/mouse genome rearrangements

Evolution preserved functional elements!



We can 'read' evolution to reveal functional elements

Today's goal:

How do we actually align two genes?

Genomes change over time

begin

A	C	G	T	C	A	T	C	A
---	---	---	---	---	---	---	---	---

A	C	G	T	G	A	T	C	A
---	---	---	---	----------	---	---	---	---

mutation

A	X	G	T	G	X	T	C	A
---	--------------	---	---	---	--------------	---	---	---

deletion

A	G	T	G	T	C	A
---	---	---	---	---	---	---

T	A	G	T	G	T	C	A
----------	---	---	---	---	---	---	---

insertion

end

T	A	G	T	G	T	C	A
---	---	---	---	---	---	---	---

Goal of alignment: Infer edit operations

begin

A	C	G	T	C	A	T	C	A
---	---	---	---	---	---	---	---	---

?




end


T	A	G	T	G	T	C	A
---	---	---	---	---	---	---	---

From Bio to CS: Formalizing the problem


- Define set of evolutionary operations (insertion, deletion, mutation)
 - Symmetric operations allow time reversibility (part of design choice)



- Define optimality criterion (min number, min cost)
 - Impossible to infer exact series of operations (Occam's razor: find min)



- Design algorithm that achieves that optimality (or approximates it)
 - Tractability of solution depends on assumptions in the formulation



Note: Not all decisions are conflicting (some are both relevant and tractable) (e.g. Pevzner vs. Sankoff and directionality in chromosomal inversions)

Formulation 1: Longest common substring

- Given two possibly related strings S1 and S2
 - What is the longest common substring? (no gaps)

S1 A C G T C A T C A

S2 T A G T G T C A



offset: +1

S1	A	C	G	T	C	A	T	C	A
	X	X	X	X	X	X			
S2		T	A	G	T	G	T	C	A

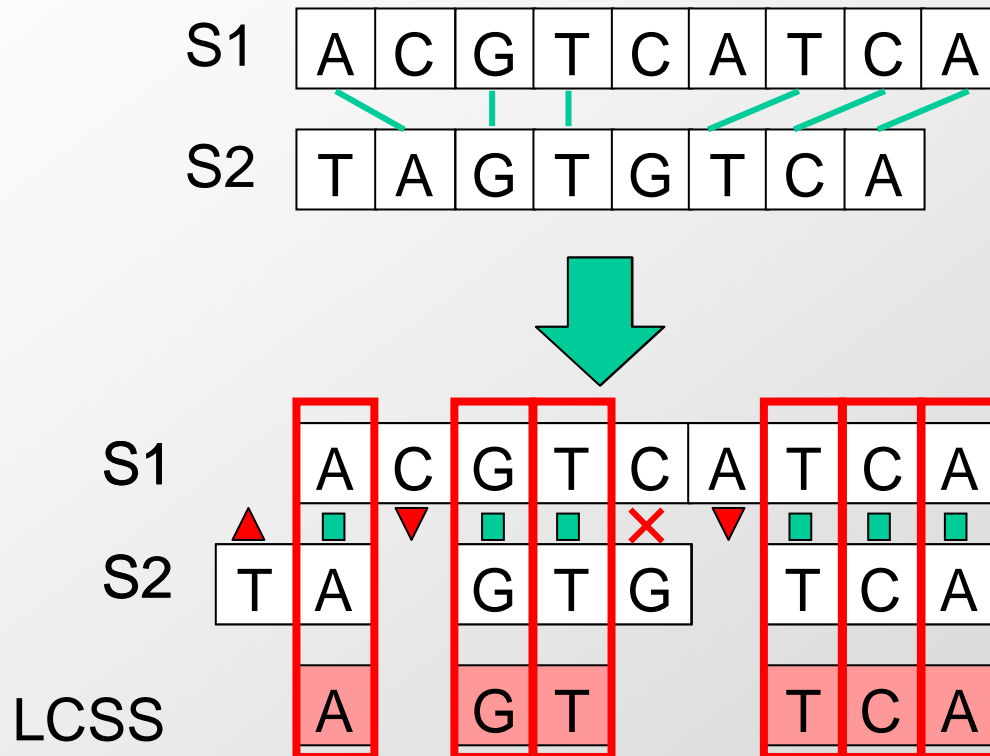


offset: -2

S1			A	C	G	T	C	A	T	C	A	
	X	X	X	X						X	X	X
S2	T	A	G	T	G	T	C	A				

Formulation 2: Longest common subsequence

- Given two possibly related strings S1 and S2
 - What is the longest common subsequence? (gaps allowed)



Edit distance:

- Number of changes needed for $S1 \rightarrow S2$
- Uniform scoring function

Formulation 3: Sequence alignment

- Allow gaps (fixed penalty)
 - Insertion & deletion operations
 - Unit cost for each character inserted or deleted
- Varying penalties for edit operations
 - Transitions (Pyrimidine \Leftrightarrow Pyrimidine, Purine \Leftrightarrow Purine)
 - Transversions (Purine \Leftrightarrow Pyrimidine changes)
 - Polymerase confuses Aw/G and Cw/T more often

Scoring function:

Match(x,x) = +1

Mismatch(A,G) = $-\frac{1}{2}$

Mismatch(C,T) = $-\frac{1}{2}$

Mismatch(x,y) = -1

	A	G	T	C
A	+1	$-\frac{1}{2}$	-1	-1
G	$-\frac{1}{2}$	+1	-1	-1
T	-1	-1	+1	$-\frac{1}{2}$
C	-1	-1	$-\frac{1}{2}$	+1

purine pyrimid.

Transitions:

$A \Leftrightarrow G$, $C \Leftrightarrow T$ common
(lower penalty)

Transversions:

All other operations

Etc...
(e.g. varying gap penalties)

How can we compute best alignment

S1	A	C	G	T	C	A	T	C	A
S2	T	A	G	T	G	T	C	A	

- Given additive scoring function:
 - Cost of mutation (AG, CT, other)
 - Cost of insertion / deletion
 - Reward of match
- Need algorithm for inferring best alignment
 - Enumeration?
 - How would you do it?
 - How many alignments are there?

Can we simply enumerate all possible alignments?


- Ways to align two sequences of length m , n

$$\binom{n+m}{m} = \frac{(m+n)!}{(m!)^2} \approx \frac{2^{m+n}}{\sqrt{\pi \cdot m}}$$


- For two sequences of length n

n	Enumeration	Today's lecture
10	184,756	100
20	1.40E+11	400
100	9.00E+58	10,000







Key insight: score is additive!



- Compute best alignment recursively
 - For a given aligned pair (i, j) , the best alignment is:
 - Best alignment of S1[1.. i] and S2[1.. j]
 - + Best alignment of S1[i ..n] and S2[j ..m]
 - Proof: cut-and-paste argument (see 6.046)



Key insight: re-use computation



Identical sub-problems! We can reuse our work!

Solution #1 – Memoization

- Create a big dictionary, indexed by aligned seqs
 - When you encounter a new pair of sequences
 - If it is in the dictionary:
 - Look up the solution
 - If it is not in the dictionary
 - Compute the solution
 - Insert the solution in the dictionary
- Ensures that there is no duplicated work
 - Only need to compute each sub-alignment once!

Top down approach

Solution #2 – Dynamic programming

- Create a big table, indexed by (i,j)
 - Fill it in from the beginning all the way till the end
 - You know that you'll need every subpart
 - Guaranteed to explore entire search space
- Ensures that there is no duplicated work
 - Only need to compute each sub-alignment once!
- Very simple computationally!

Bottom up approach

A simple introduction to Dynamic Programming

- Fibonacci numbers







Figure by MIT OpenCourseWare.




Fibonacci numbers are ubiquitous in nature




Rabbits per generation




Leaves per height




Romanesque spirals



Nautilus size



Coneflower spirals




Leaf ordering

Computing Fibonacci numbers: Top down

- Fibonacci numbers are defined recursively:
 - Python code

```
def fibonacci(n):
    if n==1 or n==2: return 1
    return fibonacci(n-1) + fibonacci(n-2)
```

- **Goal: Compute n^{th} Fibonacci number.**
 - $F(0)=1, F(1)=1, F(n)=F(n-1)+F(n-2)$
 - 1,1,2,3,5,8,13,21,34,55,89,144,233,377,...
- **Analysis:**
 - $T(n) = T(n-1) + T(n-2) = (\dots) = O(2^n)$



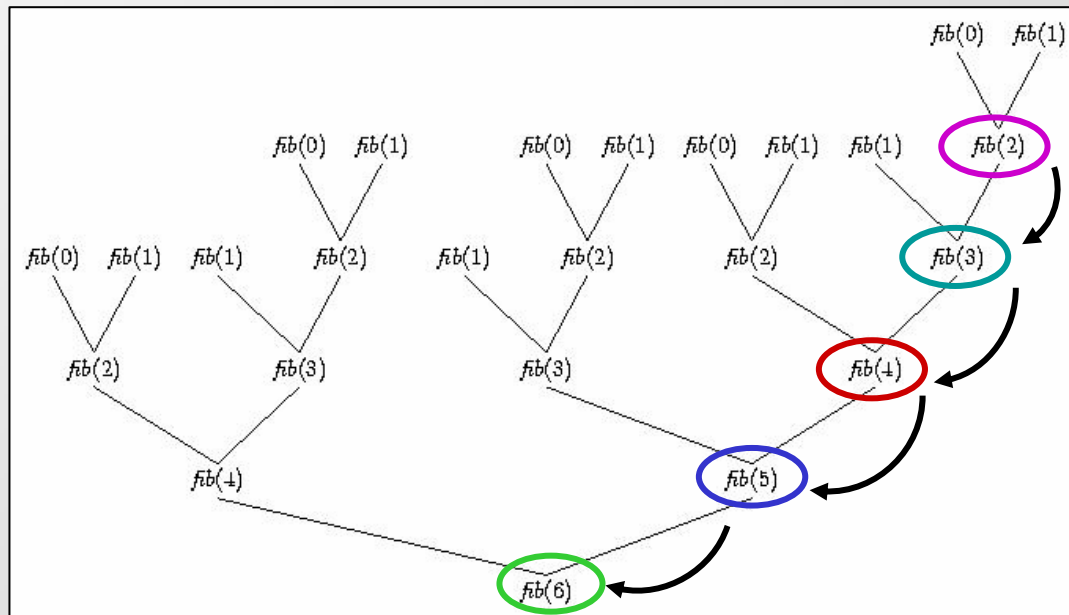
Computing Fibonacci numbers: Bottom up

- Top-down approach
 - Python code

fib_table	
F[1]	1
F[2]	1
F[3]	2
F[4]	3
F[5]	5
F[6]	8
F[7]	13
F[8]	21
F[9]	34
F[10]	55
F[11]	89
F[12]	?

```
def fibonacci(n):  
    fib_table[1] = 1  
    fib_table[2] = 1  
    for i in range(3,n+1):  
        fib_table[i] = fib_table[i-1]+fib_table[i-2]  
    return fib_table[n]
```

- Analysis: $T(n) = O(n)$



Lessons from iterative Fibonacci algorithm

fib_table	
F[1]	1
F[2]	1
F[3]	2
F[4]	3
F[5]	5
F[6]	8
F[7]	13
F[8]	21
F[9]	34
F[10]	55
F[11]	89
F[12]	?



- What did the iterative solution do?
 - Reveal identical sub-problems
 - Order computation to enable result reuse
 - Systematically filled-in table of results
 - Expressed larger problems from their subparts
- Ordering of computations matters
 - Naïve top-down approach very slow
 - results of smaller problems not available
 - repeated work
 - Systematic bottom-up approach successful
 - Systematically solve each sub-problem
 - Fill-in table of sub-problem results in order.
 - Look up solutions instead of recomputing

Dynamic Programming in Theory

- **Hallmarks of Dynamic Programming**
 - **Optimal substructure:** Optimal solution to problem (instance) contains optimal solutions to sub-problems
 - **Overlapping subproblems:** Limited number of distinct subproblems, repeated many many times
- **Typically for optimization problems (unlike Fib example)**
 - Optimal choice made locally: $\max(\text{subsolution score})$
 - Score is typically added through the search space
 - Traceback common, find optimal path from indiv. choices
- **Middle of the road in range of difficulty**
 - Easier: greedy choice possible at each step
 - DynProg: requires a traceback to find that optimal path
 - Harder: no opt. substr., e.g. subproblem dependencies

Hallmarks of optimization problems

Greedy algorithms

Dynamic Programming

1. Optimal substructure

An optimal solution to a problem (instance) contains optimal solutions to subproblems.

2. Overlapping subproblems

A recursive solution contains a “small” number of distinct subproblems repeated many times.

3. Greedy choice property

Locally optimal choices lead to globally optimal solution


*Greedy Choice is not possible
Globally optimal solution requires trace back through many choices*

Dynamic Programming in Practice


- Setting up dynamic programming
 1. Find 'matrix' parameterization (# dimensions, variables)
 2. Make sure sub-problem space is finite! (not exponential)
 - If not all subproblems are used, better off using memoization
 - If reuse not extensive, perhaps DynProg is not right solution!
 3. Traversal order: sub-results ready when you need them
 - Computation order matters! (bottom-up, but not always obvious)
 4. Recursion formula: larger problems = $F(\text{subparts})$
 5. Remember choices: typically $F()$ includes $\min()$ or $\max()$
 - Need representation for storing pointers, is this polynomial !
- Then start computing
 1. Systematically fill in table of results, find optimal score
 2. Trace-back from optimal score, find optimal solution

**How do we apply dynamic programming
to sequence alignment ?**

Key insight: score is additive!




- Compute best alignment recursively
 - For a given aligned pair (i, j) , the best alignment is:
 - Best alignment of S1[1.. i] and S2[1.. j]
 - + Best alignment of S1[i .. n] and S2[j .. m]





Dynamic Programming for sequence alignment

- Setting up dynamic programming
 1. Find 'matrix' parameterization
 2. Make sure sub-problem space is finite! (not exponential)
 3. Traversal order: sub-results ready when you need them
 4. Recursion formula: larger problems = $F(\text{subparts})$
 5. Remember choices: typically $F()$ includes $\min()$ or $\max()$
- Then start computing
 1. Systematically fill in table of results, find optimal score
 2. Trace-back from optimal score, find optimal solution

(1, 2, 3) Store score of aligning (i,j) in matrix M(i,j)



Duality: seq. alignment \Leftrightarrow path through the matrix



Goal:
Find best path
through the matrix

(4) Filling in the dynamic programming matrix

- Local update rules:
 - Compute next alignment based on previous alignment
 - Just like Fibonacci numbers: $F[i] = F[i-1] + F[i-2]$
 - Table lookup!
- Compute scores for prefixes of increasing length
 - This allows a single recursion (top-left to bottom-right) instead of two recursions (middle-to-outside top-down)
 - Only three possibilities for extending by one nucleotide: a gap in one species, a gap in the other, a (mis)match
 - When you reach bottom right, prefix of length n is seq S
- Computing the score of a cell from its neighbors
$$F(i-1, j) - \text{gap}$$
 - $F(i, j) = \max \{ \begin{array}{l} F(i, j) + \text{score} \\ F(i, j-1) - \text{gap} \end{array} \}$

0. Setting up the scoring matrix

	-	A	G	T
-	0			
A				
A				
G				
C				

Initialization:

- Top left: 0

Update Rule:

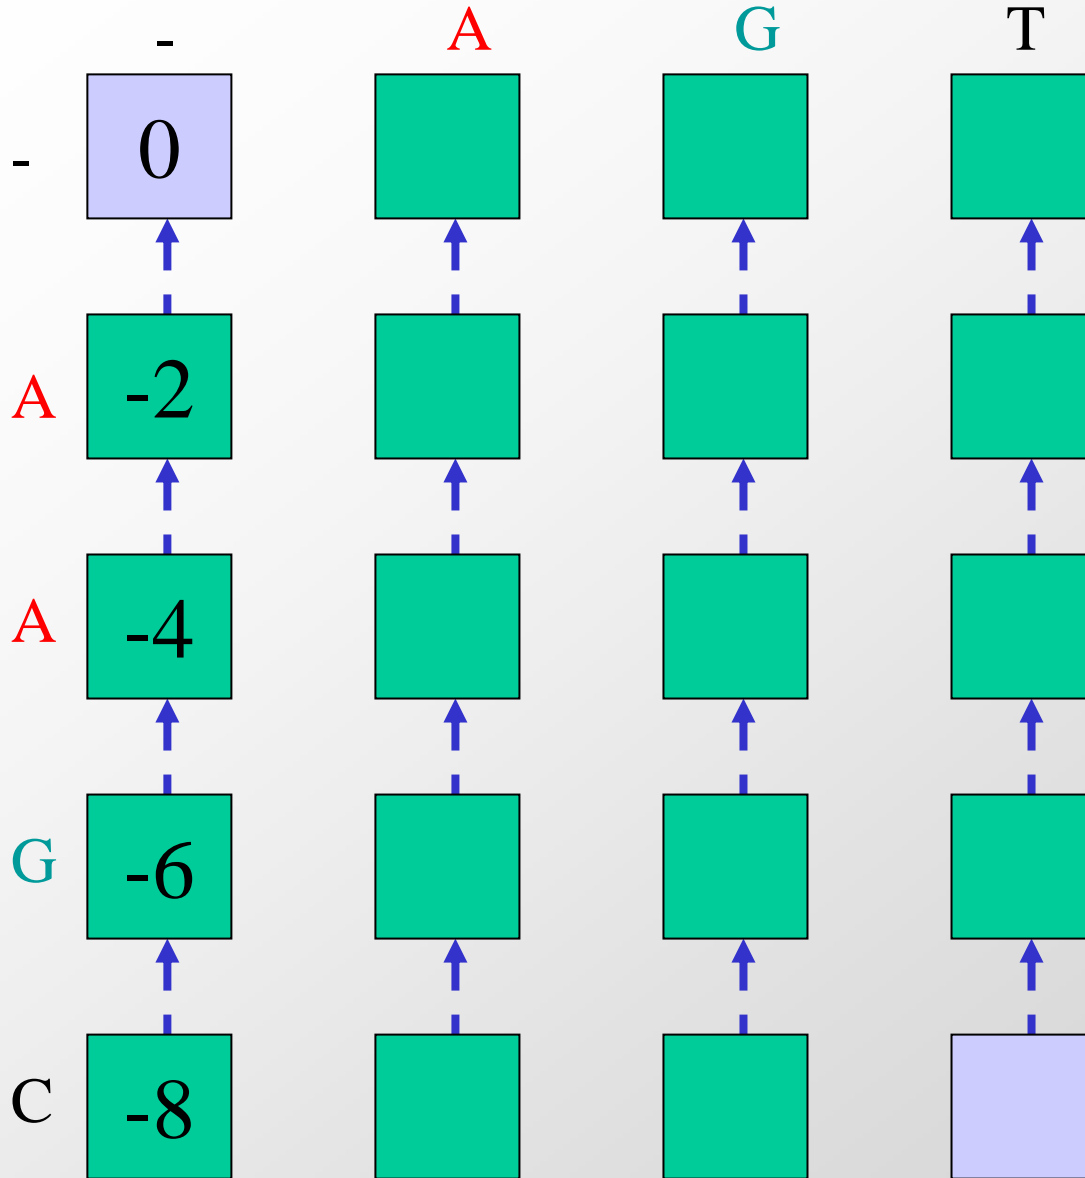
$$A(i,j)=\max\{$$

}

Termination:

- Bottom right

1. Allowing gaps in s



Initialization:

- Top left: 0

Update Rule:


$$A(i,j) = \max\{$$

- $A(i-1, j) - 2$

Termination:

- Bottom right

2. Allowing gaps in t



Initialization:

- Top left: 0

Update Rule:


$$A(i,j) = \max\{$$

- $A(i-1, j) - 2$
- $A(i, j-1) - 2$

Termination:

- Bottom right

3. Allowing mismatches



Initialization:

- Top left: 0

Update Rule:


$A(i,j) = \max\{$

- $A(i-1, j) - 2$
- $A(i, j-1) - 2$
- $A(i-1, j-1) - 1$

Termination:

- Bottom right

4. Choosing optimal paths



Initialization:

- Top left: 0

Update Rule:


$$A(i,j) = \max\{$$

- $A(i-1, j) - 2$
- $A(i, j-1) - 2$
- $A(i-1, j-1) - 1$

Termination:

- Bottom right

5. Rewarding matches



Initialization:

- Top left: 0

Update Rule:

$A(i,j) = \max\{$


- $A(i-1, j) - 2$
- $A(i, j-1) - 2$
- $A(i-1, j-1) \pm 1$

Termination:

- Bottom right

What is missing? (5) Returning the actual path!

- We know how to compute the best score
 - Simply the number at the bottom right entry
- But we need to remember where it came from
 - Pointer to the choice we made at each step
- Retrace path through the matrix
 - Need to remember all the pointers




Time needed: $O(m*n)$

Space needed: $O(m*n)$

Summary

- Dynamic programming
 - Reuse of computation
 - Order sub-problems. Fill table of sub-problem results
 - Read table instead of repeating work (ex: Fibonacci)
- Sequence alignment
 - Edit distance and scoring functions
 - Dynamic programming matrix
 - Matrix traversal path \Leftrightarrow Optimal alignment
- Thursday: Variations on sequence alignment
 - Local/global alignment, affine gaps, algo speed-ups
 - Semi-numerical alignment, hashing, database lookup
- Recitation:
 - Dynamic programming applications
 - Probabilistic derivations of alignment scores

Bounded Dynamic Programming



Initialization:

$F(i,0), F(0,j)$ undefined for $i, j > k$

Iteration:

For $i = 1 \dots M$


For $j = \max(1, i - k) \dots \min(N, i + k)$

$$F(i, j) = \max \begin{cases} F(i - 1, j - 1) + s(x_i, y_j) \\ F(i, j - 1) - d, \text{ if } j > i - k(N) \\ F(i - 1, j) - d, \text{ if } j < i + k(N) \end{cases}$$

Termination: same

Linear space alignment

It is easy to compute $F(M, N)$ in linear space



Allocate (column[1])

Allocate (column[2])

For $i = 1 \dots M$

If $i > 1$, then:

Free(column[$i - 2$])

Allocate(column[i])

For $j = 1 \dots N$



$F(i, j) = \dots$

What about the pointers?

Finding the best back-pointer for current column


- Now, using 2 columns of space, we can compute for $k = 1 \dots M$, $F(M/2, k)$, $F^r(M/2, N-k)$

PLUS the backpointers




Best forward-pointer for current column

- Now, we can find k^* maximizing $F(M/2, k) + F^r(M/2, N-k)$
- Also, we can trace the path exiting column $M/2$ from k^*




Recursively find midpoint for left & right

- Iterate this procedure to the left and right!



Total time cost of linear-space alignment



Total Time: $cMN + cMN/2 + cMN/4 + \dots = 2cMN = O(MN)$

Total Space: $O(N)$ for computation,
 $O(N+M)$ to store the optimal alignment

Formulation 4: Varying gap cost models (next time)

(still) Varying penalties for edit operations

Now allow gaps of varying penalty:

1. Linear gap penalty
 - Same as before,
2. Affine gap penalty
 - Big initial cost for starting or ending a gap
 - Small incremental cost for each additional character
3. General gap penalty
 - Any cost function
 - No longer computable using the same model
4. Seek duplicated regions, rearrangements, ...