

MIT OpenCourseWare  
<http://ocw.mit.edu>

6.047 / 6.878 Computational Biology: Genomes, Networks, Evolution  
Fall 2008

For information about citing these materials or our Terms of Use, visit: <http://ocw.mit.edu/terms>.

# **Rapid sequence alignment and Database search**

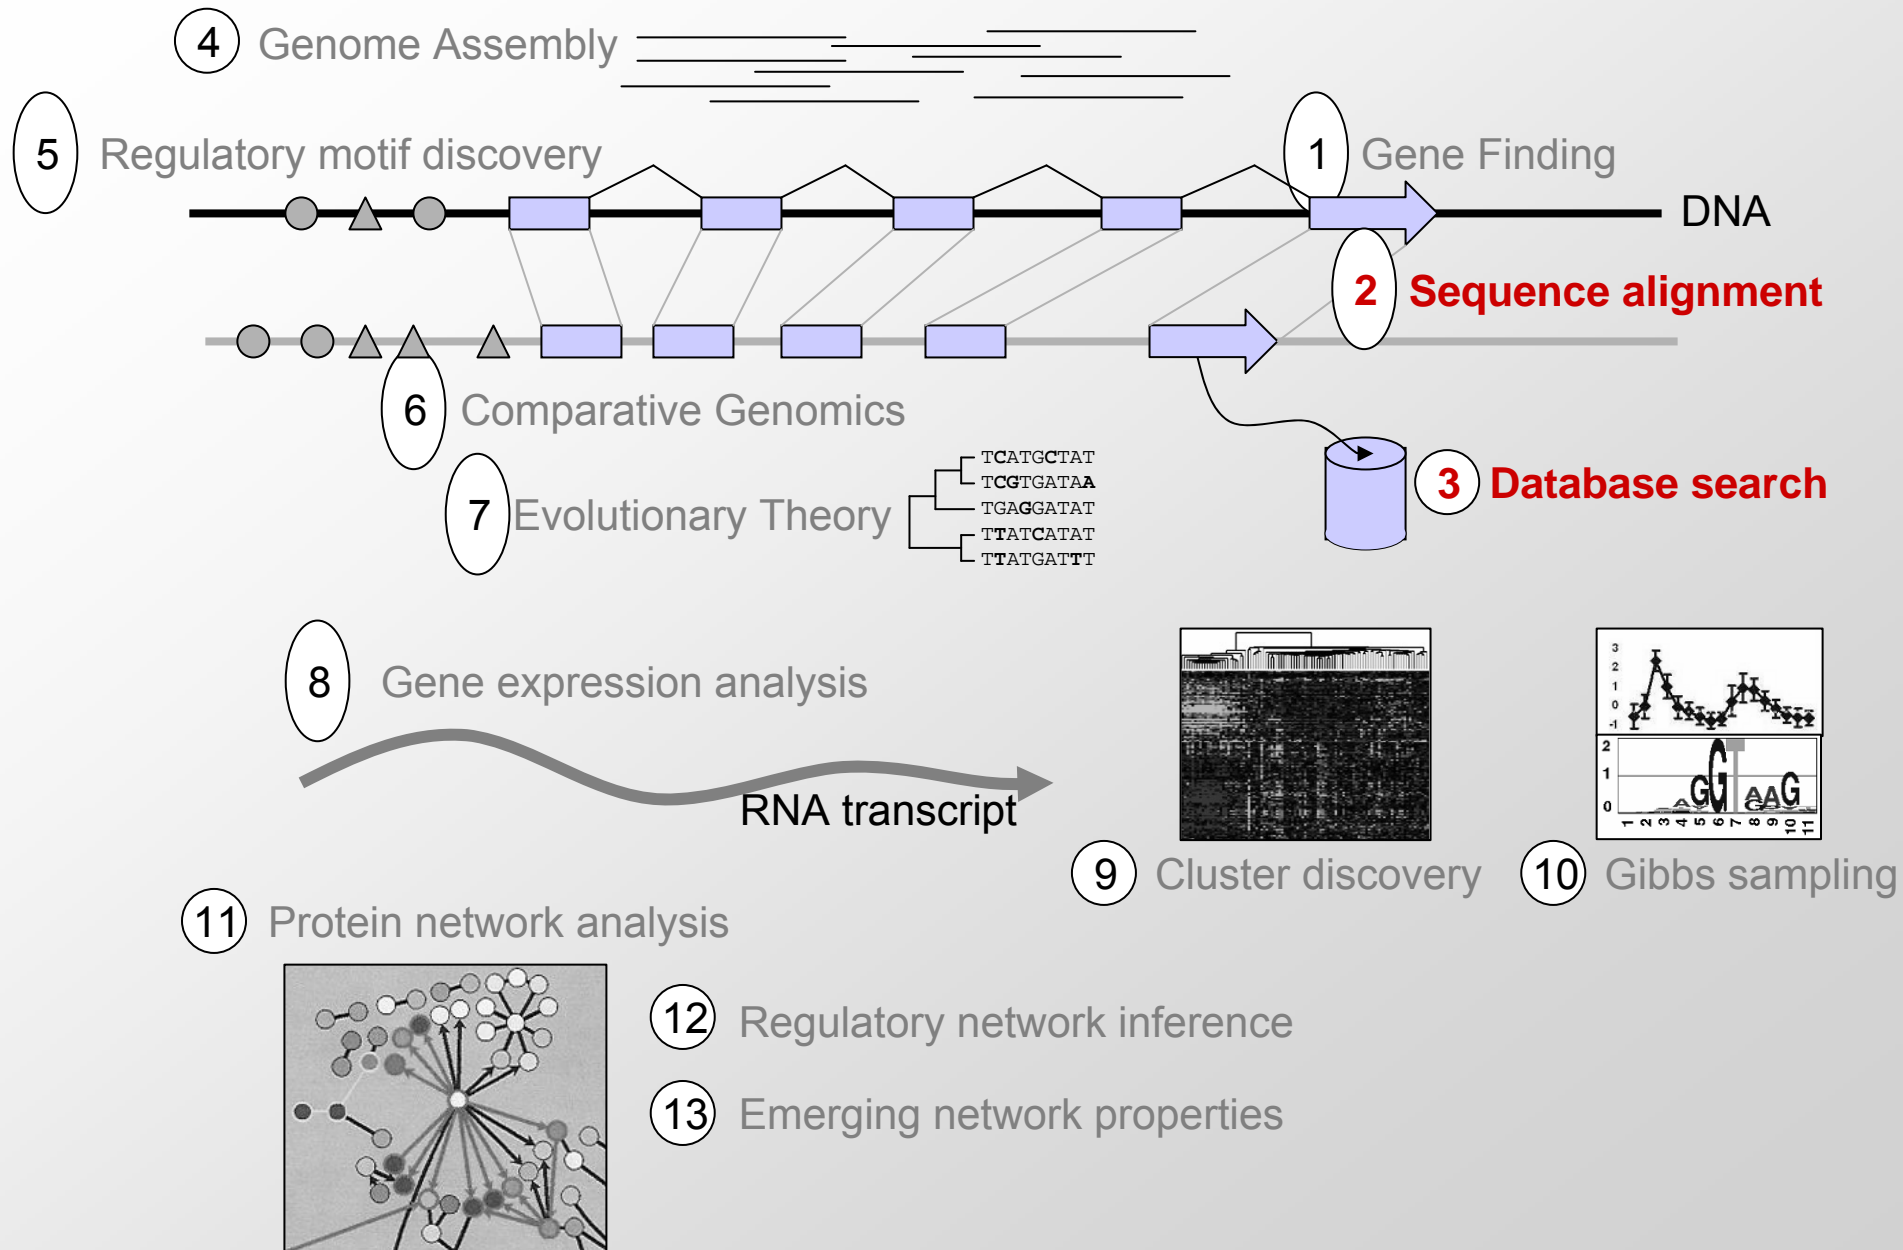
Local alignment, varying gap penalties

Karp-Rabin: Semi-numerical methods


BLAST: dB search, neighborhood search

Statistics of alignment scores (recitation)


# Challenges in Computational Biology




# Tues: Sequence alignment + dynamic programming



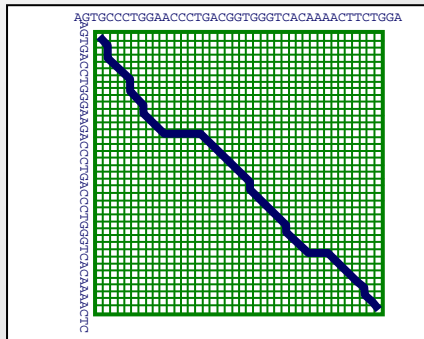
- Dynamic programming
  - Problems that can be decomposed into subparts
  - Identical sub-problems: reuse computation
  - Bottom-up approach: systematically fill table



- The sequence alignment problem
  - Genomes change: mutation, insertions, deletions
  - Alignment: infer evolutionary events
  - Scoring metric reflects evolutionary properties



- Dynamic programming and sequence alignment
  - Alignment scores are additive: decomposable
  - Represent sub-problem scores in  $M(i,j)$  matrix
  - Duality between alignment and path through matrix



- Needleman-Wunsch algorithm
  - Local update rule:  $F(i,j) = \max\{\text{up}, \text{left}, \text{diagonal}\}$
  - Save choice pointers for traceback
  - Bottom-right corner gives optimal alignment score
  - Trace-back of pointers gives optimal path/alignment

# Today's Goal: Diving deeper into alignments

## 1. Global alignment vs. Local alignment

- Needleman-Wunsch and Smith-Waterman
- Varying gap penalties and algorithmic speedups

## 2. Linear-time exact string matching

- Karp-Rabin algorithm and semi-numerical methods
- Hash functions and randomized algorithms

## 3. The BLAST algorithm and inexact matching

- Hashing with neighborhood search
- Two-hit blast and hashing with combs

## 4. Probabilistic foundations of sequence alignment

- Mismatch penalties, BLOSUM and PAM matrices
- Statistical significance of an alignment score

# Today's Goal: Diving deeper into alignments

## 1. Global alignment vs. Local alignment

- Needleman-Wunsch and Smith-Waterman
- Varying gap penalties and algorithmic speedups

## 2. Linear-time exact string matching

- Karp-Rabin algorithm and semi-numerical methods
- Hash functions and randomized algorithms

## 3. The BLAST algorithm and inexact matching



- Hashing with neighborhood search
- Two-hit blast and hashing with combs

## 4. Probabilistic foundations of sequence alignment


- Mismatch penalties, BLOSUM and PAM matrices
- Statistical significance of an alignment score

# Intro to Local Alignments


- Statement of the problem
  - A *local alignment* of strings  $s$  and  $t$  is an alignment of a substring of  $s$  with a substring of  $t$
- Why local alignments?
  - Small domains of a gene may be only conserved portions
  - Looking for a small gene in a large chromosome (search)
  - Large segments often undergo rearrangements



Global alignment





Local alignment



# Global Alignment

vs.

# Local alignment



## Needleman-Wunsch algorithm

Initialization:  $F(0, 0) = 0$

Iteration:

$$F(i, j) = \max \begin{cases} F(i-1, j) - d \\ F(i, j-1) - d \\ F(i-1, j-1) + s(x_i, y_j) \end{cases}$$

Termination: Bottom right

## Smith-Waterman algorithm

Initialization:  $F(0, j) = F(i, 0) = 0$

Iteration:


$$F(i, j) = \max \begin{cases} 0 \\ F(i-1, j) - d \\ F(i, j-1) - d \\ F(i-1, j-1) + s(x_i, y_j) \end{cases}$$

Termination: Anywhere




# More variations on the theme: semi-global alignment


- Sequence alignment variations



Global



Local



Semi-global

Initialization

Top left

Top row/left col.

Top row

Iteration: max

$$F(i-1, j) - d$$

$$F(i, j-1) - d$$

$$F(i-1, j-1) + s(x_i, y_j)$$

$$0$$

$$F(i-1, j) - d$$

$$F(i, j-1) - d$$

$$F(i-1, j-1) + s(x_i, y_j)$$

$$F(i-1, j) - d$$

$$F(i, j-1) - d$$

$$F(i-1, j-1) + s(x_i, y_j)$$


Termination

Bottom right

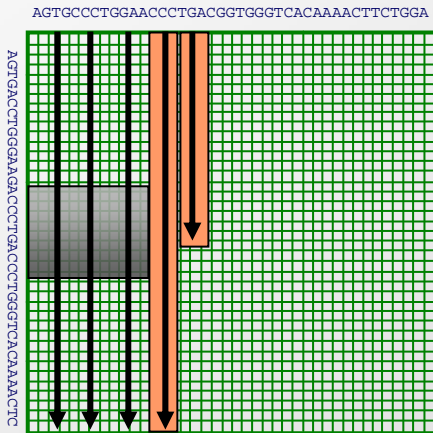
Anywhere

Right column

# Some algorithmic variations to save time/space




- Save time: Bounded-space computation
  - Space:  $O(k*m)$
  - Time:  $O(k*m)$ , where  $k$  = radius explored
  - Heuristic
    - Not guaranteed optimal answer
    - Works very well in practice
  - Practical interest



- Save space: Linear-space computation
  - Save only one col / row / diag at a time
  - Computes optimal score easily
  - Recursive call modification allows traceback
  - Theoretical interest
    - Effective running time slower
    - Optimal answer guaranteed

# Sequence alignment with generalized gap penalties

- Implementing a generalized gap penalty function  $F(\text{gap\_length})$



Initialization: same

Iteration:

$$F(i, j) = \max \begin{cases} F(i-1, j-1) + s(x_i, y_j) \\ \max_{k=0 \dots i-1} F(k, j) - \gamma(i-k) \\ \max_{k=0 \dots j-1} F(i, k) - \gamma(j-k) \end{cases}$$


Termination: same

Running Time:  $O(N^2M)$  (cubic)

Space:  $O(NM)$


Do we have to be  
so general?

# Algorithmic trade-offs of varying gap penalty functions




Linear gap penalty:  $w(k) = k * p$

- State: Current index tells if in a gap or not
- Achievable using quadratic algorithm (even w/ linear space)




Quadratic:  $w(k) = p + q * k + r k^2$ .

- State: needs to encode the length of the gap, which can be  $O(n)$
- To encode it we need  $O(\log n)$  bits of information. Not feasible



Affine gap penalty:  $w(k) = p + q * k$ , where  $q < p$

- State: add binary value for each sequence: starting a gap or not
- Implementation: add second matrix for already-in-gap (recitation)



Length (mod 3) gap penalty for protein-coding regions

- Gaps of length divisible by 3 are penalized less: conserve frame
- This is feasible, but requires more possible states
- Possible states are: starting, mod 3=1, mod 3=2, mod 3=0

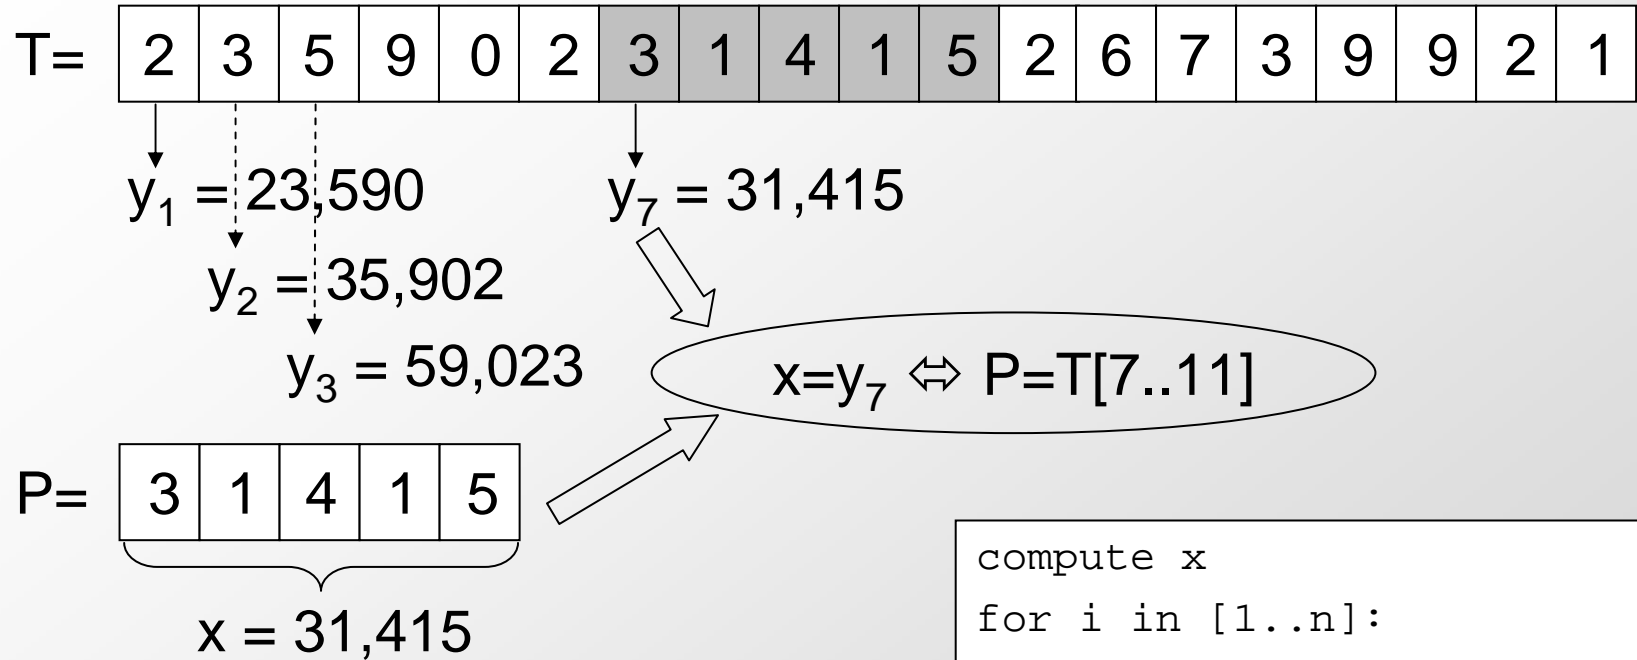
# Today's Goal: Diving deeper into alignments

1. Global alignment vs. Local alignment
  - Needleman-Wunsch and Smith-Waterman
  - Varying gap penalties and algorithmic speedups
- 2. Linear-time exact string matching**
  - Karp-Rabin algorithm and semi-numerical methods
  - Hash functions and randomized algorithms
3. The BLAST algorithm and inexact matching
  - Hashing with neighborhood search
  - Two-hit blast and hashing with combs
4. Probabilistic foundations of sequence alignment
  - Mismatch penalties, BLOSUM and PAM matrices
  - Statistical significance of an alignment score

# Linear-time string matching

- When looking for exact matches of a pattern
- Karp-Rabin algorithm: interpret it numerically
  - Start with ‘broken’ version of the algorithm
  - Progressively fix it to make it work
- Several other solutions exist, not covered today:
  - Z-algorithm / fundamental pre-processing, Gusfield
  - Boyer-Moore and Knuth-Morris-Pratt algorithms are earliest instantiations, similar in spirit
  - Suffix trees: beautiful algorithms, many different variations and applications, limited use in CompBio
  - Suffix arrays: practical variation, Gene Myers

# Karp-Rabin algorithm

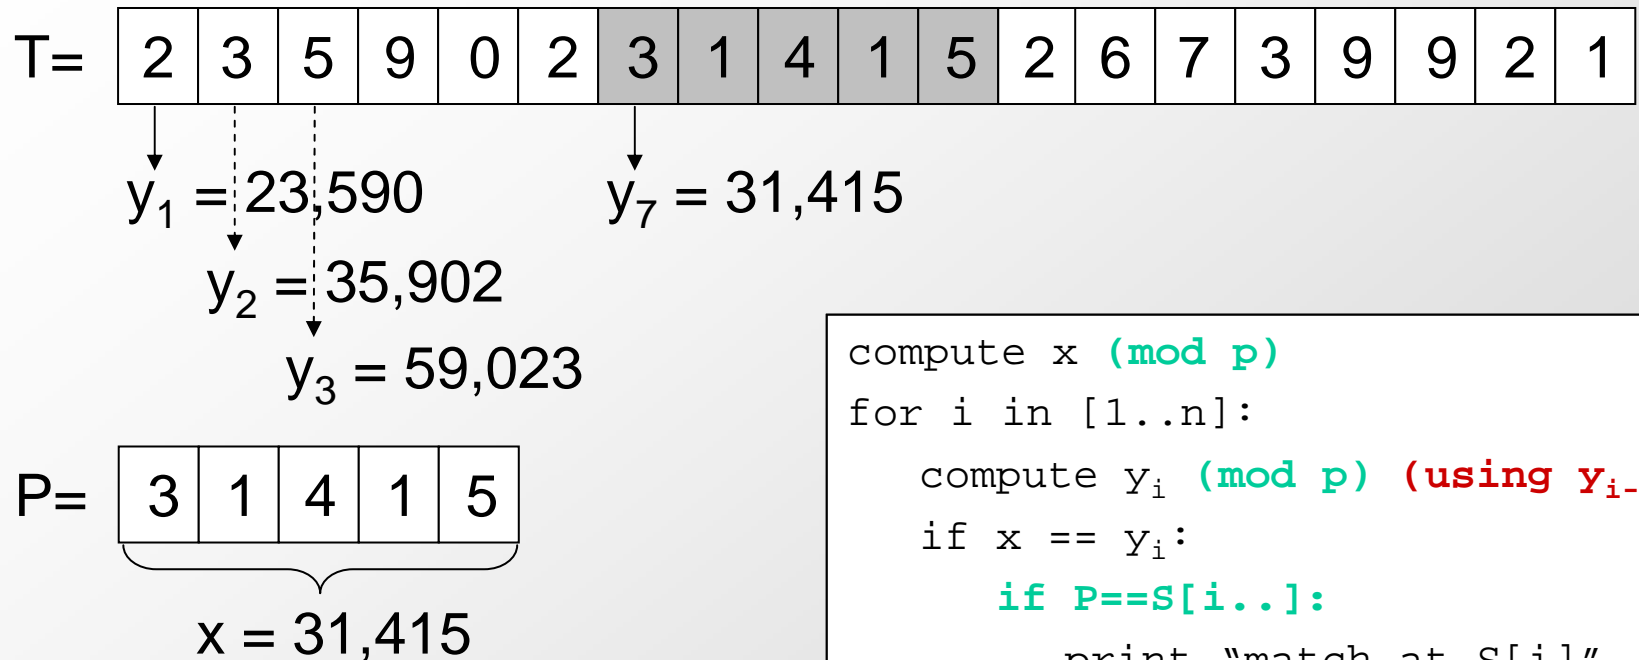


```
compute x
for i in [1..n]:
    compute yi
    if x == yi:
        print "match at S[i]"
```

**(this does not actually work)**

- Key idea:
  - Interpret strings as numbers: fast comparison

# Karp-Rabin algorithm





```
compute x (mod p)
for i in [1..n]:
    compute  $y_i$  (mod p) (using  $y_{i-1}$ )
    if x ==  $y_i$ :
        if P==S[i..]:
            print "match at S[i]"
        else:
            (spurious hit)
```

(this actually works)

- Key idea:
  - Interpret strings as numbers: fast comparison
- To make it work:
  - Compute next number based on previous one  $\rightarrow O(1)$
  - Hashing (mod p)  $\rightarrow$  keep the numbers small  $\rightarrow O(1)$



# Hashing is good, but leads to collisions



- Consequences of (mod p) 'hashing'
    - Good: Enable fast computation (use small numbers)
    - Bad: Leads to spurious hits (collisions)
- ➔ Complete algorithm must deal with the bad

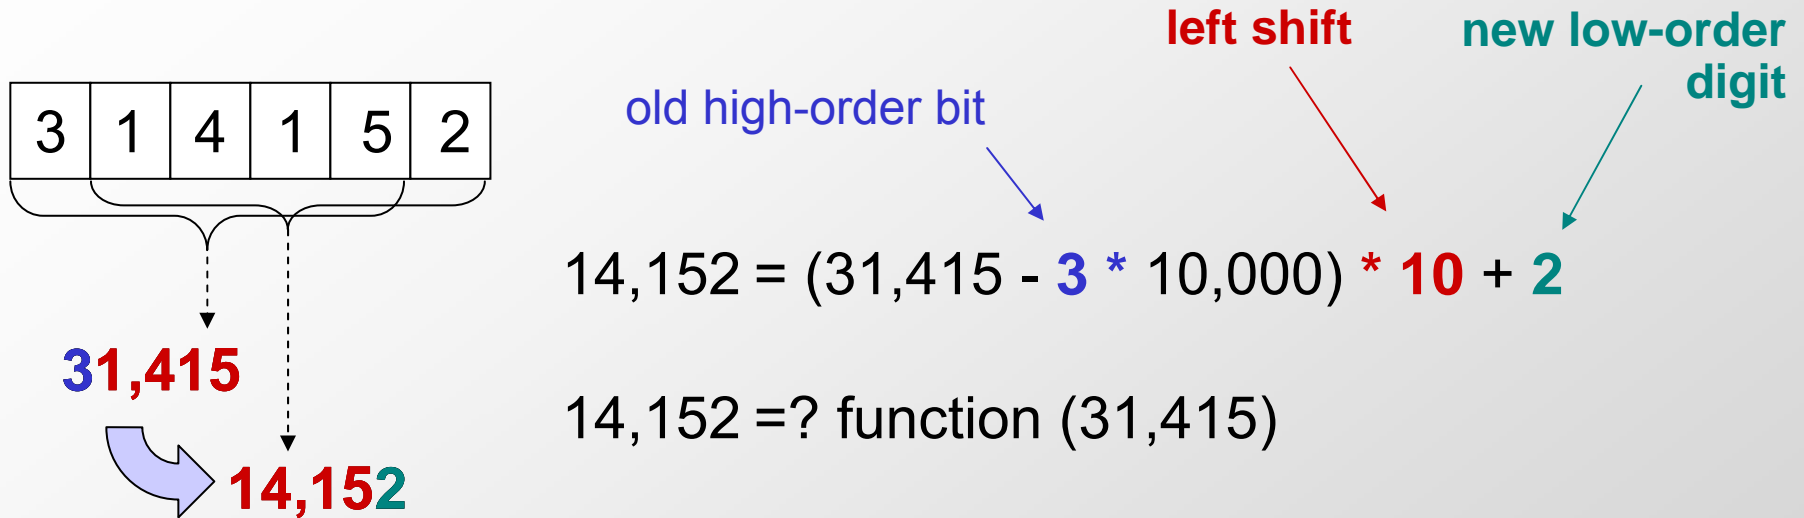
# Karp Rabin key idea: Semi-numerical approach

- Idea 1: *semi-numerical* approach:
  - Consider all  $m$ -mers:  
 $T[1\dots m], T[2\dots m+1], \dots, T[m-n+1\dots n]$
  - Map each  $T[s+1\dots s+m]$  into a *number*  $t_s$
  - Map the pattern  $P[1\dots m]$  into a *number*  $p$
  - Report the  $m$ -mers that map to the same value as  $p$

# Semi-numerical approach: implementation

- First attempt:
  - Assume  $\Sigma=\{0,1\}$   
(for  $\{A,G,T,C\}$  convert:  $A \rightarrow 00$ ,  $G \rightarrow 01$ ,  $A \rightarrow 10$ ,  $G \rightarrow 11$ )
  - Think about each  $T[s+1\dots s+m]$  as a number in binary representation, i.e.,
$$t_s = T[s+1]2^{m-1} + T[s+2]2^{m-2} + \dots + T[s+m]2^0$$
  - Output all  $s$  such that  $t_s$  is equal to the number  $p$  represented by  $P$
- Problem: how to map all  $m$ -mers in  $O(n)$  time ?
  - Find a fast way of computing  $t_{s+1}$  given  $t_s$

# Computing $t_{s+1}$ based on $t_s$ in constant time



- Middle digits of the number are already computed  
Shift them to the left ←
- Remove the high-order bit
- Add the low-order bit

## Idea 2: Computing all numbers in linear time

- How to transform

$$t_s = \underline{T[s+1]2^{m-1}} + T[s+2]2^{m-2} + \dots + T[s+m]2^0$$

Into


$$t_{s+1} = T[s+2]2^{m-1} + T[s+3]2^{m-2} + \dots + \underline{T[s+m+1]2^0} ?$$

- Can compute  $t_{s+1}$  from  $t_s$  using 3 arithmetic operations:
  - Subtract  $T[s+1]2^{m-1}$
  - Multiply by 2 (i.e., shift the bits by one position)
  - Add  $T[s+m+1]2^0$
- Therefore:  $t_{s+1} = (t_s - T[s+1]2^{m-1}) * 2 + T[s+m+1]2^0$
- Therefore, we can compute all  $t_0, t_1, \dots, t_{n-m}$  using  $O(n)$  arithmetic operations, and a number for  $P$  in  $O(m)$

## Problem: Long strings = big numbers

- To get  $O(n)$  time, we would need to perform each arithmetic operation in  $O(1)$  time
- However, the arguments are  $m$ -bit long !
- If  $m$  large, it is unreasonable to assume that operations on such big numbers can be done in  $O(1)$  time
- We need to reduce the number range to something more manageable

# Dealing with long numbers in constant time



old high-order bit

shift

new low-order digit

$$\begin{aligned} 14,152 &= (31,415 - \text{old high-order bit} * 10,000) * \text{shift} + \text{new low-order digit} \pmod{13} \\ &= (7 - 3 * 3) * 10 + 2 \pmod{13} \\ &= 8 \pmod{13} \end{aligned}$$

## Idea 3: Hashing

- We will instead compute

$$t'_s = T[s+1]2^{m-1} + T[s+2]2^{m-2} + \dots + T[s+m]2^0 \bmod q$$

where  $q$  is an “appropriate” prime number

- One can still compute  $t'_{s+1}$  from  $t'_s$  :

$$t'_{s+1} = (t'_s - T[s+1]2^{m-1}) * 2 + T[s+m+1]2^0 \bmod q$$

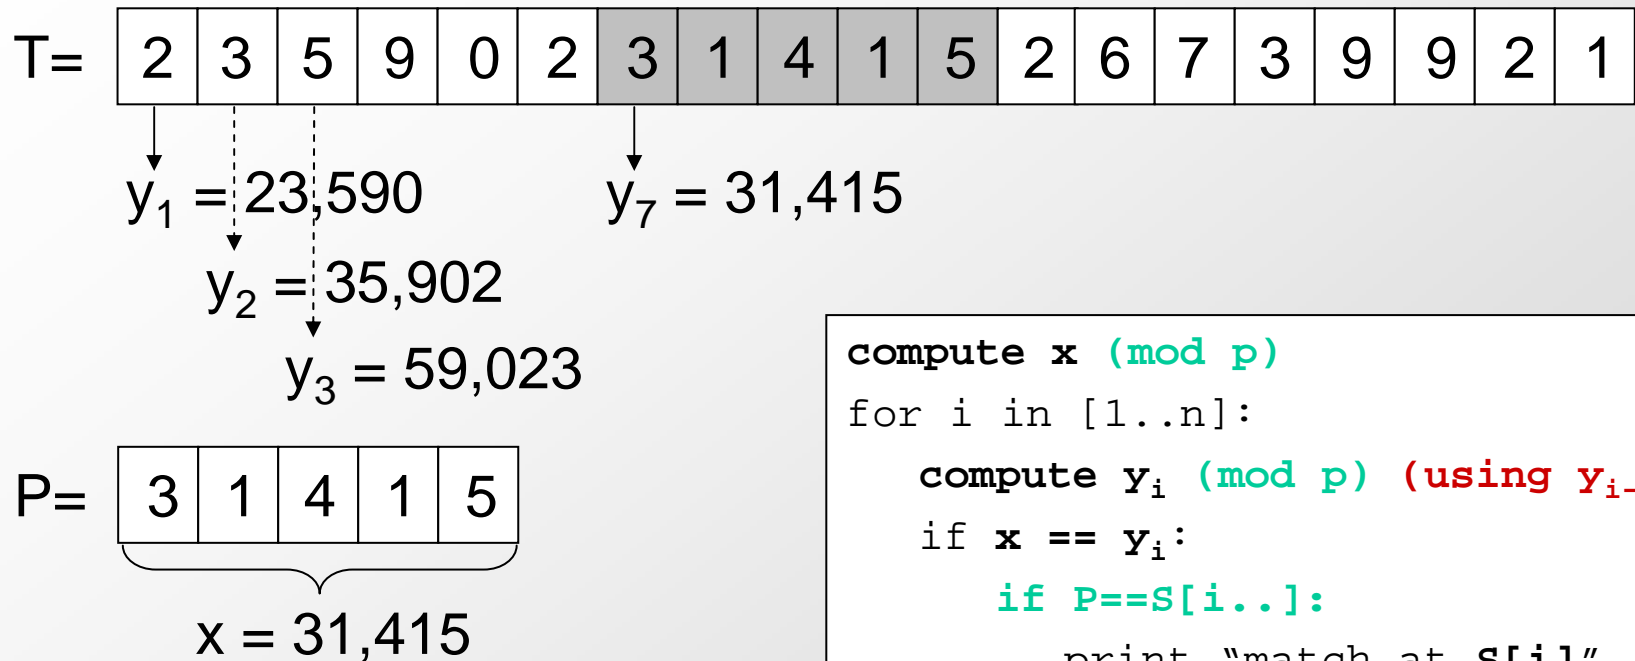
- If  $q$  is not large, we can compute all  $t'_s$  (and  $p'$ ) in  $O(n)$  time



## Problem: hashing leads to false positives

- Unfortunately, we can have false positives, i.e.,  $T[s+1 \dots s+m] \neq P$  but  $t_s \bmod q = p \bmod q$
- Our approach:
  - Use a random  $q$
  - Show that the probability of a false positive is small  
→ randomized algorithm

# Karp-Rabin algorithm: Putting it all together



```
compute x (mod p)
for i in [1..n]:
    compute  $y_i$  (mod p) (using  $y_{i-1}$ )
    if  $x == y_i$ :
        if P==S[i..]:
            print "match at S[i]"
        else:
            (spurious hit)
```

(this actually works)

- Key idea: Semi-numerical computation
  - Idea 1: Interpret strings as numbers => fast comparison  
(other semi-numerical methods: Fast Fourier Transform, Shift-And)
- To make it work:
  - Idea 2: Compute next number based on previous one  $\rightarrow O(1)$
  - Idea 3: Hashing (mod p)  $\rightarrow$  keep the numbers small  $\rightarrow O(1)$

# Today's Goal: Diving deeper into alignments

1. Global alignment vs. Local alignment
  - Needleman-Wunsch and Smith-Waterman
  - Varying gap penalties and algorithmic speedups
2. Linear-time exact string matching
  - Karp-Rabin algorithm and semi-numerical methods
  - Hash functions and randomized algorithms
- 3. The BLAST algorithm and inexact matching**
  - Hashing with neighborhood search
  - Two-hit blast and hashing with combs
4. Probabilistic foundations of sequence alignment
  - Mismatch penalties, BLOSUM and PAM matrices
  - Statistical significance of an alignment score

# Increased sequence availability → new problems

- Global Alignment and Dyn. Prog. Applications
  - Assume sequences have some common ancestry
  - Finding the “right” alignment between two sequences
    - Find minimum number of transformation operations
  - Understanding evolutionary events: mutations, indels
- Sequence databases
  - Query: new sequence. Subject: many old sequences
  - Goal: which sequences are related to the one at hand
  - most sequences will be completely unrelated to query
  - Individual alignment needs not be perfect.
    - Once initial matches are reported, can fine-tune them later
  - Query must be very fast for a new sequence

# Speeding up your searches

- Exploit nature of the problem
  - If you're going to reject any match with  $\text{idperc} \leq 90$ , then why bother even looking at sequences which don't have a stretch of 10 nucleotides in a row.
  - Pre-screen sequences for common long stretches
- Put the speed where you need it
  - Pre-processing the database is off-line.
  - Once the query arrives, must act fast
- Solution: content-based indexing and BLAST
  - Example: index 10-mers.
  - Only one 10-mer in  $4^{10}$  will match, one in a million.
  - (even with 500 k-mers, only 1 in 2000 will match).
  - Additional speedups...

# BLAST

**Basic local alignment search tool** - all 46 versions »

SF Altschul, W Gish, W Miller, EW Myers, DJ Lipman - J. Mol. Biol, 1990

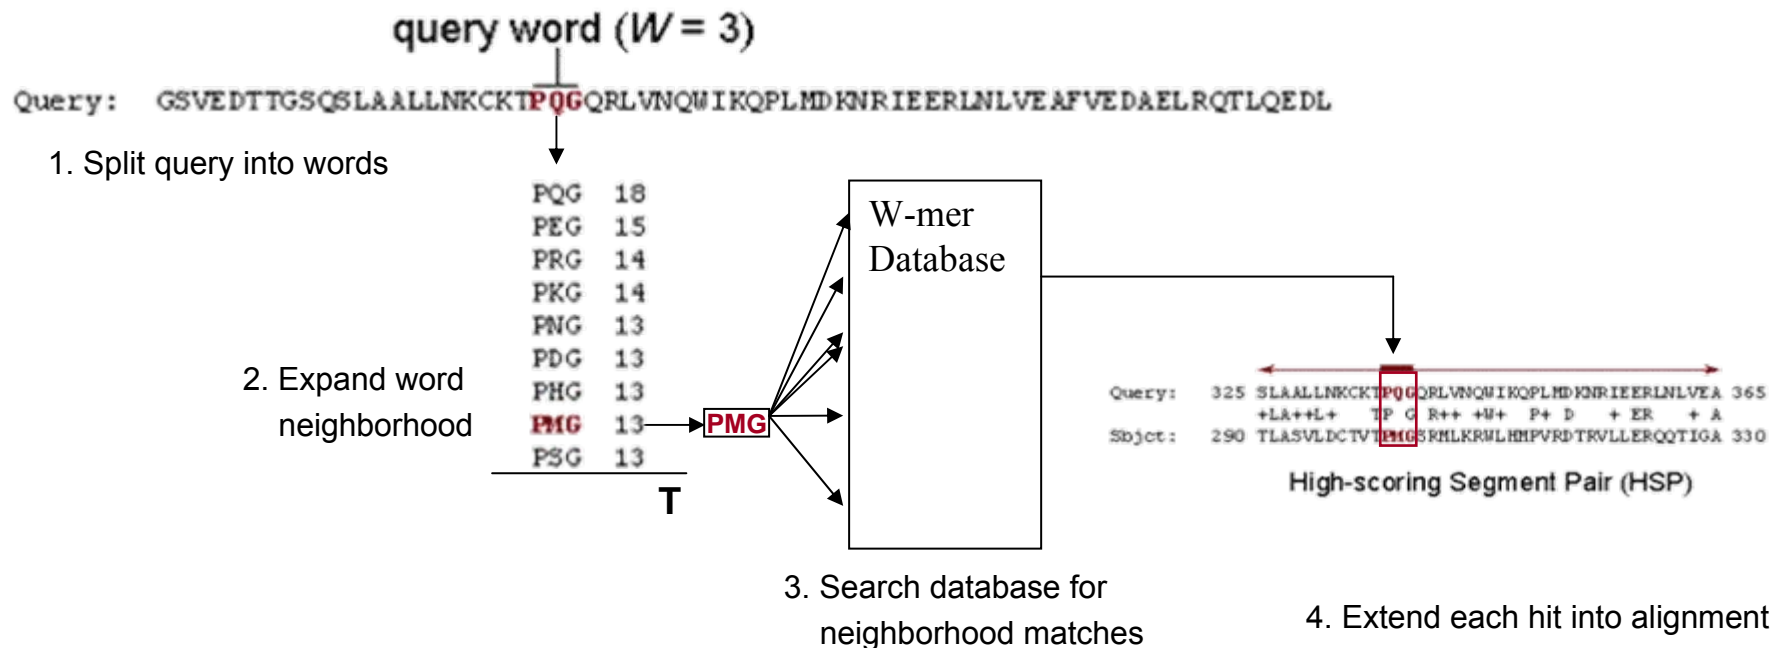
Gish', Webb Miller<sup>2</sup> Eugene W. Myers<sup>3</sup> and David J. Lipman<sup>1</sup> ...

Cited by **21457** - Related Articles - View as HTML - Web Search

(Gapped blast: 24000 citations!)

# Blast Algorithm Overview

- Receive query
  1. Split query into overlapping words of length  $W$
  2. Find neighborhood words for each word until threshold  $T$
  3. Look into the table where these neighbor words occur: seeds  $S$
  4. Extend seeds  $S$  until score drops off under  $X$
- Report significance and alignment of each match



# The BLAST Search Algorithm

query word ( $W = 3$ )

Query: GSVEDTTGSQSLAALLNKCKT**PQG**QRLVNQWIKQPLMDKNRIEERLNLVEAFVEDAELRQTLQEDL

neighborhood  
words

PQG	18
PEG	15
PRG	14
PKG	14
PNG	13
PDG	13
PHG	13
<b>PMG</b>	13
PSG	13
PQA	12
PQN	12
etc. ...	

neighborhood  
score threshold  
( $T = 13$ )

**X**

←-----→


Query: 325 SLAALLNKCKT**PQ**GQRLVNQWIKQPLMDKNRIEERLNLVEA 365  
           +LA++L+ TP G R++ +W+ P+ D + ER + A

Sbjct: 290 TLASVLDCTVT**PMG**SRMLKRULHMPVRDTRVLLERQQTIGA 330

High-scoring Segment Pair (HSP)




# Why BLAST works(1): Pigeonhole and W-mers



- Pigeonhole principle

- If you have 2 pigeons and 3 holes, there must be at least one hole with no pigeon



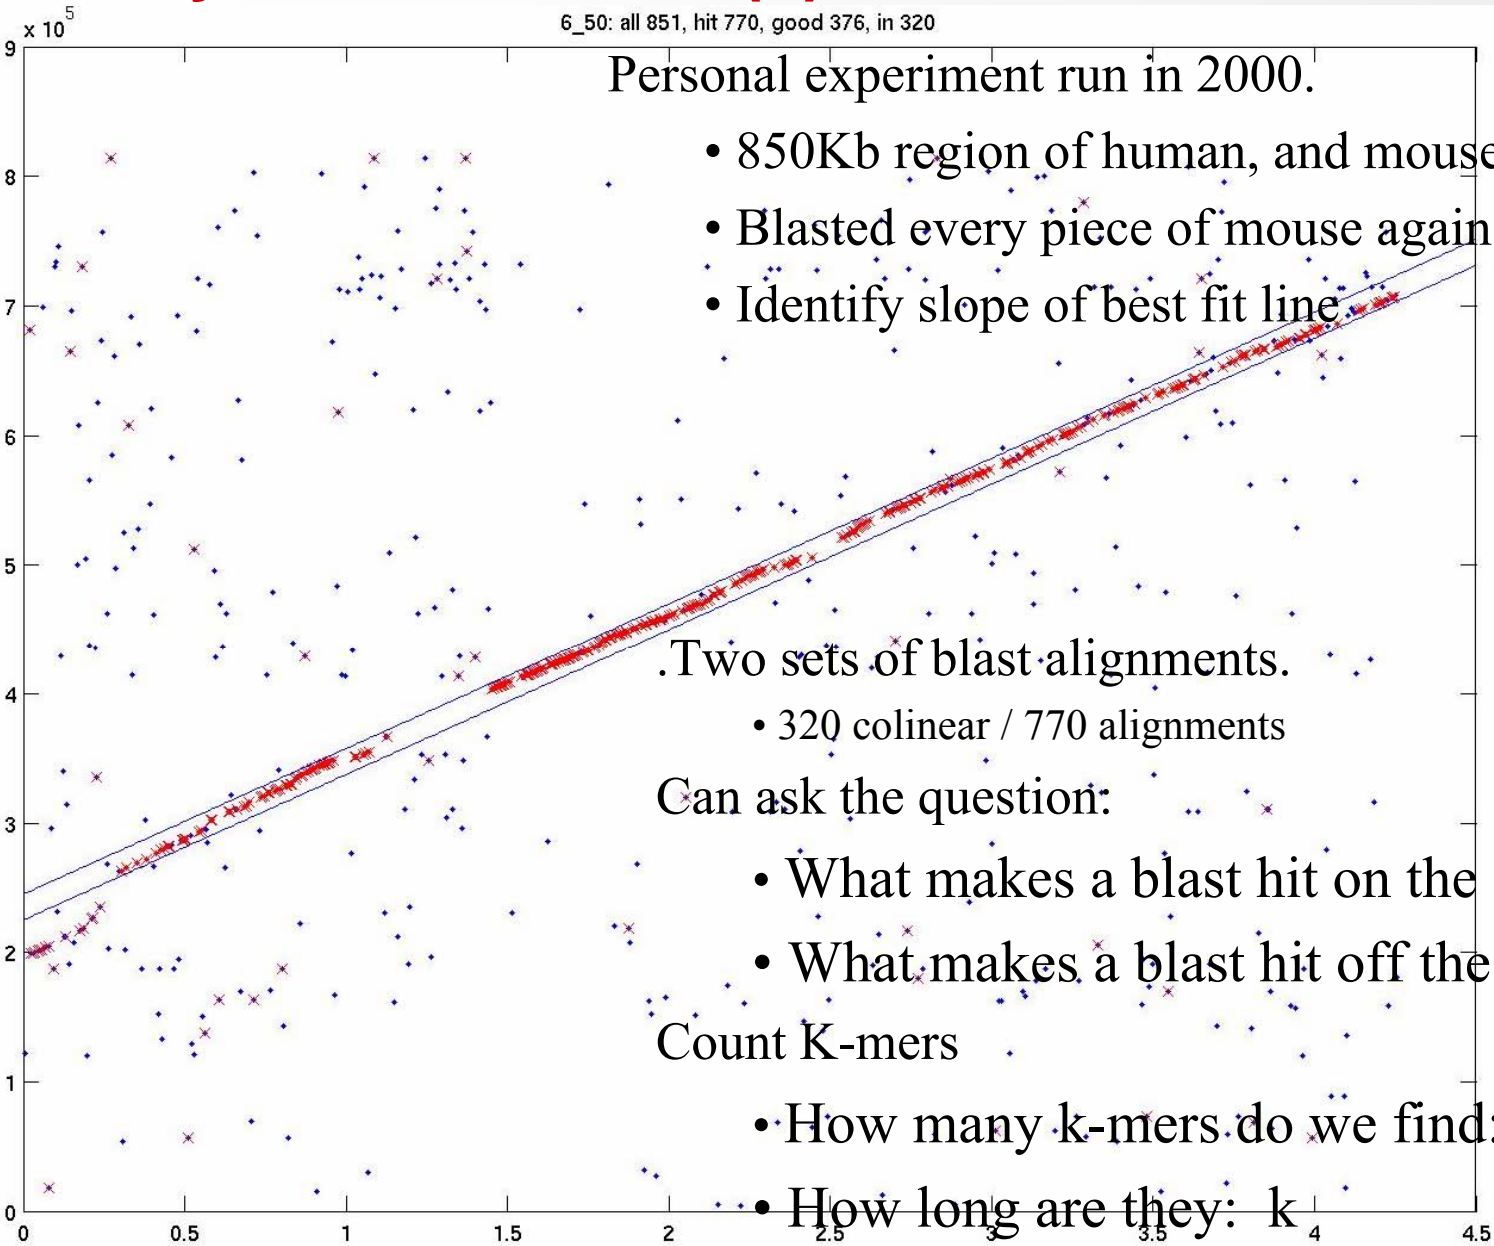
- Pigeonholing mis-matches

- Two sequences, each 9 amino-acids, with 7 identities
- There is a stretch of 3 amino-acids perfectly conserved

- In general:

- Sequence length:  $n$
- Identities:  $t$
- Can use  $W$ -mers for  $W = \lceil n/(n-t+1) \rceil$

# Why BLAST works(2): K-mer matches in practice



- 850Kb region of human, and mouse 450Kb ortholog.
- Blasted every piece of mouse against human (6,50)
- Identify slope of best fit line.

.Two sets of blast alignments.

- 320 colinear / 770 alignments


Can ask the question:

- What makes a blast hit on the line look good.
- What makes a blast hit off the diagonal look bad.

- How many k-mers do we find:  $n$
- How long are they:  $k$

# True alignments: Looking for K-mers

number of k-mers that happen for each length of k-mer.



# Extensions to the basic algorithm

- Ideas beyond W-mer indexing ?
  - Faster
  - Better sensitivity (less false negatives)
- 1. Filtering: Low complexity regions cause spurious hits
  - Filter out low complexity in your query
  - Filter most over-represented items in your database
- 2. Two-hit BLAST
  - Two smaller W-mers are more likely than one longer one
  - Therefore it's a more sensitive searching method to look for two hits instead of one, with the same speed.
  - Improves sensitivity for any speed, speed for any sensitivity
- 3. Beyond W-mers, hashing with Combs

## Extension(3): Combs and Random Projections

### Key idea:

- No reason to use only consecutive symbols
- Instead, we could use **combs**, e.g.,  
RGIKW  $\rightarrow$  R\*IK\* , RG\*\*W, ...
- Indexing same as for W-mers:
  - For each comb, store the list of positions in the database where it occurs
  - Perform lookups to answer the query
- How to choose the combs? At random
  - Randomized projection:  
Califano-Rigoutsos'93, Buhler'01, Indyk-Motwani'98
  - Choose the positions of \* at random
  - Analyze false positives and false negatives

# Extension(3): Combs and Random Projections

## Performance Analysis:

- Assume we select  $k$  positions, which do **not** contain \*, at random **with replacement**
- What is the probability of a false negative ?
  - At most:  $1 - \text{idperc}^k$
  - In our case:  $1 - (7/9)^4 = 0.63...$
- What is we repeat the process  $l$  times, independently ?
  - Miss prob. =  $0.63^l$
  - For  $l=5$ , it is less than 10%

Query: RKIWGDPRS

Datab.: RKI**V**GDR**R**S

$k=4$



Query: \*KI\*G\*\*\*S

Datab.: \*KI\*G\*\*\*S



# Today's Goal: Diving deeper into alignments

## 1. Global alignment vs. Local alignment

- Needleman-Wunsch and Smith-Waterman
- Varying gap penalties and algorithmic speedups

## 2. Linear-time exact string matching

- Karp-Rabin algorithm and semi-numerical methods
- Hash functions and randomized algorithms

## 3. The BLAST algorithm and inexact matching

- Hashing with neighborhood search
- Two-hit blast and hashing with combs

## 4. Probabilistic foundations of sequence alignment

- Mismatch penalties, BLOSUM and PAM matrices
- Statistical significance of an alignment score

# Varying scores/penalties for matches/mismatches

Nucleotide sequences

	A	G	T	C
A	+1	-1/2	-1	-1
G	-1/2	+1	-1	-1
T	-1	-1	+1	-1/2
C	-1	-1	-1/2	+1

purine      pyrimid.

**Transitions:**

$A \leftrightarrow G$ ,  $C \leftrightarrow T$  common  
(lower penalty)

**Transversions:**

All other operations

Protein space: amino-acid similarities

	C	S	T	P	A	G	N	D	E	Q	H	R	K	M	I	L	V	F	Y	W	
C	9																				C
S	-1	4																			S
T	-1	1	5																		T
P	-3	-1	-1	7																	P
A	0	1	0	-1	4																A
G	-3	0	-2	-2	0	6															G
N	-3	1	0	-2	-2	0	6														N
D	-3	0	-1	-1	-2	-1	1	6													D
E	-4	0	-1	-1	-1	-2	0	2	5												E
Q	-3	0	-1	-1	-1	-2	0	0	2	5											Q
H	-3	-1	-2	-2	-2	-2	1	-1	0	0	8										H
R	-3	-1	-1	-2	-1	-2	0	-2	0	1	0	5									R
K	-3	0	-1	-1	-1	-2	0	-1	1	1	-1	2	5								K
M	-1	-1	-1	-2	-1	-3	-2	-3	-2	0	-2	-1	-1	5							M
I	-1	-2	-1	-3	-1	-4	-3	-3	-3	-3	-3	-3	-3	1	4						I
L	-1	-2	-1	-3	-1	-4	-3	-4	-3	-2	-3	-2	-2	2	2	4					L
V	-1	-2	0	-2	0	-3	-3	-3	-2	-2	-3	-3	-2	1	3	1	4				V
F	-2	-2	-2	-4	-2	-3	-3	-3	-3	-3	-1	-3	-3	0	0	0	-1	6			F
Y	-2	-2	-2	-3	-2	-3	-2	-3	-2	-1	2	-2	-2	-1	-1	-1	-1	3	7		Y
W	-2	-3	-2	-4	-3	-2	-4	-4	-3	-2	-2	-3	-3	-1	-3	-2	-3	1	2	11	W

**BLOSUM matrix of AA similarity scores**

- Where do these scores come from?
- Are two aligned sequences actually related?



# Probabilistic Model of Alignments

- we'll focus on protein alignments without gaps
- given an alignment, we can consider two possibilities
  - R**: the sequences are related by evolution
  - U**: the sequences are unrelated
- How can we distinguish these possibilities?
- How is this view related to amino-acid substitution matrices?

# Model for Unrelated Sequences

- we'll assume that each position in the alignment is sampled randomly from some distribution of amino acids
- let  $q_a$  be the probability of amino acid  $a$
- the probability of an  $n$ -character alignment of  $x$  and  $y$  is given by

$$\Pr(x, y \mid U) = \prod_{i=1}^n q_{x_i} \prod_{i=1}^n q_{y_i}$$

# Model for Related Sequences

- we'll assume that each pair of aligned amino acids evolved from a common ancestor
- let  $p_{ab}$  be the probability that evolution gave rise to amino acid  $a$  in one sequence and  $b$  in another sequence
- the probability of an alignment of  $x$  and  $y$  is given by

$$\Pr(x, y \mid R) = \prod_{i=1}^n p_{x_i y_i}$$

# Probabilistic Model of Alignments

- How can we decide which possibility ( $U$  or  $R$ ) is more likely?
- one principled way is to consider the relative likelihood of the two possibilities (the odds ratio)

$$\frac{\Pr(x, y \mid R)}{\Pr(x, y \mid U)} = \frac{\prod_i p_{x_i y_i}}{\prod_i q_{x_i} \prod_i q_{y_i}} = \frac{\prod_i p_{x_i y_i}}{\prod_i q_{x_i} q_{y_i}}$$

- taking the log, we get

$$\log \frac{\Pr(x, y \mid R)}{\Pr(x, y \mid U)} = \sum_i \log \left( \frac{p_{x_i y_i}}{q_{x_i} q_{y_i}} \right)$$

# Probabilistic Model of Alignments

- the score for an alignment is thus given by:

$$S = \sum_i s(x_i, y_i) = \log \frac{\Pr(x, y \mid R)}{\Pr(x, y \mid U)}$$

- the substitution matrix score for the pair  $a, b$  should thus be given by:

$$s(a, b) = \log \left( \frac{p_{ab}}{q_a q_b} \right)$$

# Substitution Matrices

- two popular sets of matrices for protein sequences
  - PAM matrices [Dayhoff *et al.*, 1978]
  - BLOSUM matrices  
[Henikoff & Henikoff, 1992]
- both try to capture the the relative substitutability of amino acid pairs in the context of evolution

# BLOSUM62

Positive for chemically similar substitution

Common amino acids have low weights

Rare amino acids have high weights

# BLOSUM62

A	4																					
R	-1	5																				
N	-2	0	6																			
D	-2	-2	1	6																		
C	0	-3	-3	-3	9																	
Q	-1	1	0	0	-3	5																
E	-1	0	0	2	-4	2	5															
G	0	-2	0	-1	-3	-2	-2	6														
H	-2	0	1	-1	-3	0	0	-2	8													
I	-1	-3	-3	-3	-1	-3	-3	-4	-3	4												
L	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4											
K	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5										
M	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5									
F	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6								
P	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7							
S	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4						
T	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5					
W	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11				
Y	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7			
V	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4		
X	0	-1	-1	-1	-2	-1	-1	-1	-1	-1	-1	-1	-1	-1	-2	0	0	-2	-1	-1	-1	
	A	R	N	D	C	Q	E	G	H	I	L	K	M	F	P	S	T	W	Y	V	X	

Positive for chemically similar substitution

Common amino acids have low weights

Rare amino acids have high weights

# Substitution Matrices

- the substitution matrix score for the pair  $a, b$  is given by:

$$s(a, b) = \log \left( \frac{p_{ab}}{q_a q_b} \right)$$

- but how do we get values for  $p_{ab}$  (probability that  $a$  and  $b$  arose from a common ancestor)?
- it depends on how long ago sequences diverged
  - diverged recently:  $p_{ab} \approx 0$  for  $a \neq b$
  - diverged long ago:  $p_{ab} \approx q_a q_b$



# Substitution Matrices


- key idea: trusted alignments of related sequences provide information about biologically permissible mutations

# BLOSUM Matrices

- [Henikoff & Henikoff, *PNAS* 1992]
- probabilities estimated from “blocks” of sequence fragments that represent structurally conserved regions in proteins
- transition frequencies observed directly by identifying blocks that are at least
  - 45% identical (BLOSUM-45)
  - 50% identical (BLOSUM-50)
  - 62% identical (BLOSUM-62)
  - etc.

# BLOSUM Matrices

- given: a set of sequences in a block
- fill in matrix  $A$  with number of observed substitutions (we won't worry about details of some normalization that happens here)



$$p_{ab} = \frac{A_{ab}}{\sum_{c,d} A_{cd}}$$

$$q_a = \frac{\sum_b A_{ab}}{\sum_{c,d} A_{cd}}$$

**(you are not responsible for the  
remainder of this section)**

# Statistics of Alignment Scores

**Q:** How do we assess whether an alignment provides good evidence for homology?

**A:** determine how likely it is that such an alignment score would result from chance.

3 ways to calculate chance; look at alignment scores for

- real but non-homologous sequences
- real sequences shuffled to preserve compositional properties
- sequences generated randomly based upon a DNA/protein sequence model

# Statistics of Alignment Scores

- earlier we considered how do decide if a single alignment was more likely due to relatedness or chance
- but what if we're considered many alignments?
  - e.g. what if we're doing a BLAST search against a large protein database?
- we'd like to know how many high-scoring alignments we're likely to get by chance

# Distribution of Scores

- Karlin & Altschul, *PNAS*, 1990
- consider a random model in which
  - we're looking for ungapped local alignments
  - the lengths of the sequences in each pair are  $m$  and  $n$
- the expected number of alignments,  $E$ , with score at least  $S$  is given by:

$$E(S) = Kmne^{-\lambda S}$$

# Distribution of Scores

$$E(S) = Kmne^{-\lambda S}$$

- $S$  is a given score threshold
- $m$  and  $n$  are the lengths of the sequences under consideration
- $K$  and  $\lambda$  are constants that can be calculated from
  - the substitution matrix
  - the frequencies of the individual amino acids

$K$  = measure of the relative independence of points in context of MSP score

$\lambda$  = the unique positive-valued solution to  $\sum_{i,j} S_{i,j} P_x(i) P_y(j) e^{\lambda S_{ij}} = 1$



# Statistics of Alignment Scores

- to generalize this to searching a database, have  $n$  represent the summed length of the sequences in the DB
- the NCBI BLAST server does just this
- with this analysis, can also calculate  $p$ -values (the probability of a random alignment scoring at least  $S$ )
- theory for gapped alignments not as well developed
- computational experiments suggest this analysis holds for gapped alignments (but  $K$  and  $\lambda$  must be estimated from data)

# Summary: Diving deeper into sequence alignment

## 1. Global alignment vs. Local alignment

- Needleman-Wunsch and Smith-Waterman
- Varying gap penalties and algorithmic speedups

## 2. Linear-time exact string matching

- Karp-Rabin algorithm and semi-numerical methods
- Hash functions and randomized algorithms

## 3. The BLAST algorithm and inexact matching

- Hashing with neighborhood search
- Two-hit blast and hashing with combs

## 4. Probabilistic foundations of sequence alignment

- Mismatch penalties, BLOSUM and PAM matrices
- Statistical significance of an alignment score

# Tomorrow's recitation: Deeper into Alignments


- Affine gap penalties
  - Augmenting the state-space
  - Linear, affine, piecewise linear, general gap penalty
- Statistical significance of alignment
  - Where does  $s(\mathbf{x}_i, \mathbf{y}_j)$  come from?
  - Are two aligned sequences actually related

## **3c. Massive pre-processing**


Suffix Trees

# Suffix trees

- Great tool for text processing
  - E.g., searching for exact occurrence of a pattern
- Suffix tree for: xabxac




# Suffix tree definition



- **Definition:** Suffix tree  $ST$  for text  $T[1..n]$ 
  - Rooted, directed tree  $T$ ,  $n$  leaves, numbered  $1..n$
  - Text labels on the edges
  - Path to leaf  $i$  spells out the suffix  $S[i..]$ , by concatenating edge labels
  - Common prefixes share common paths, diverge to form internal nodes

# Properties of suffix trees




- How much space do we need to represent a suffix tree of  $T[1..n]$  ?
- Only  $O(n)$ 
  - At most  $O(n)$  edges
  - Each edge label can be represented as  $T[i..j]$

# Exact string matching with suffix trees

- Given the suffix tree for text  $T$
- Search for pattern  $P[1..m]$ 
  - For every character in  $P$ , traverse the appropriate path of the tree, reading one character each time
  - If  $P$  is not found in a path,  $P$  does not occur in  $T$
  - If  $P$  is found in its entirety, then all occurrences of  $P$  in  $T$  are exactly the children of that node
    - Every child corresponds to exactly one occurrence
    - Simply list each of the leaf indices
- Time:  $O(m)$


T: xabxac

P: abx





# Suffix Tree Construction



- Running time:  $O(n^2)$
- Can be improved to  $O(n)$